

# Hybrid cluster with kubernetes

Maxim Vorobjov  
Solutions Architect



**comae**  
technologies

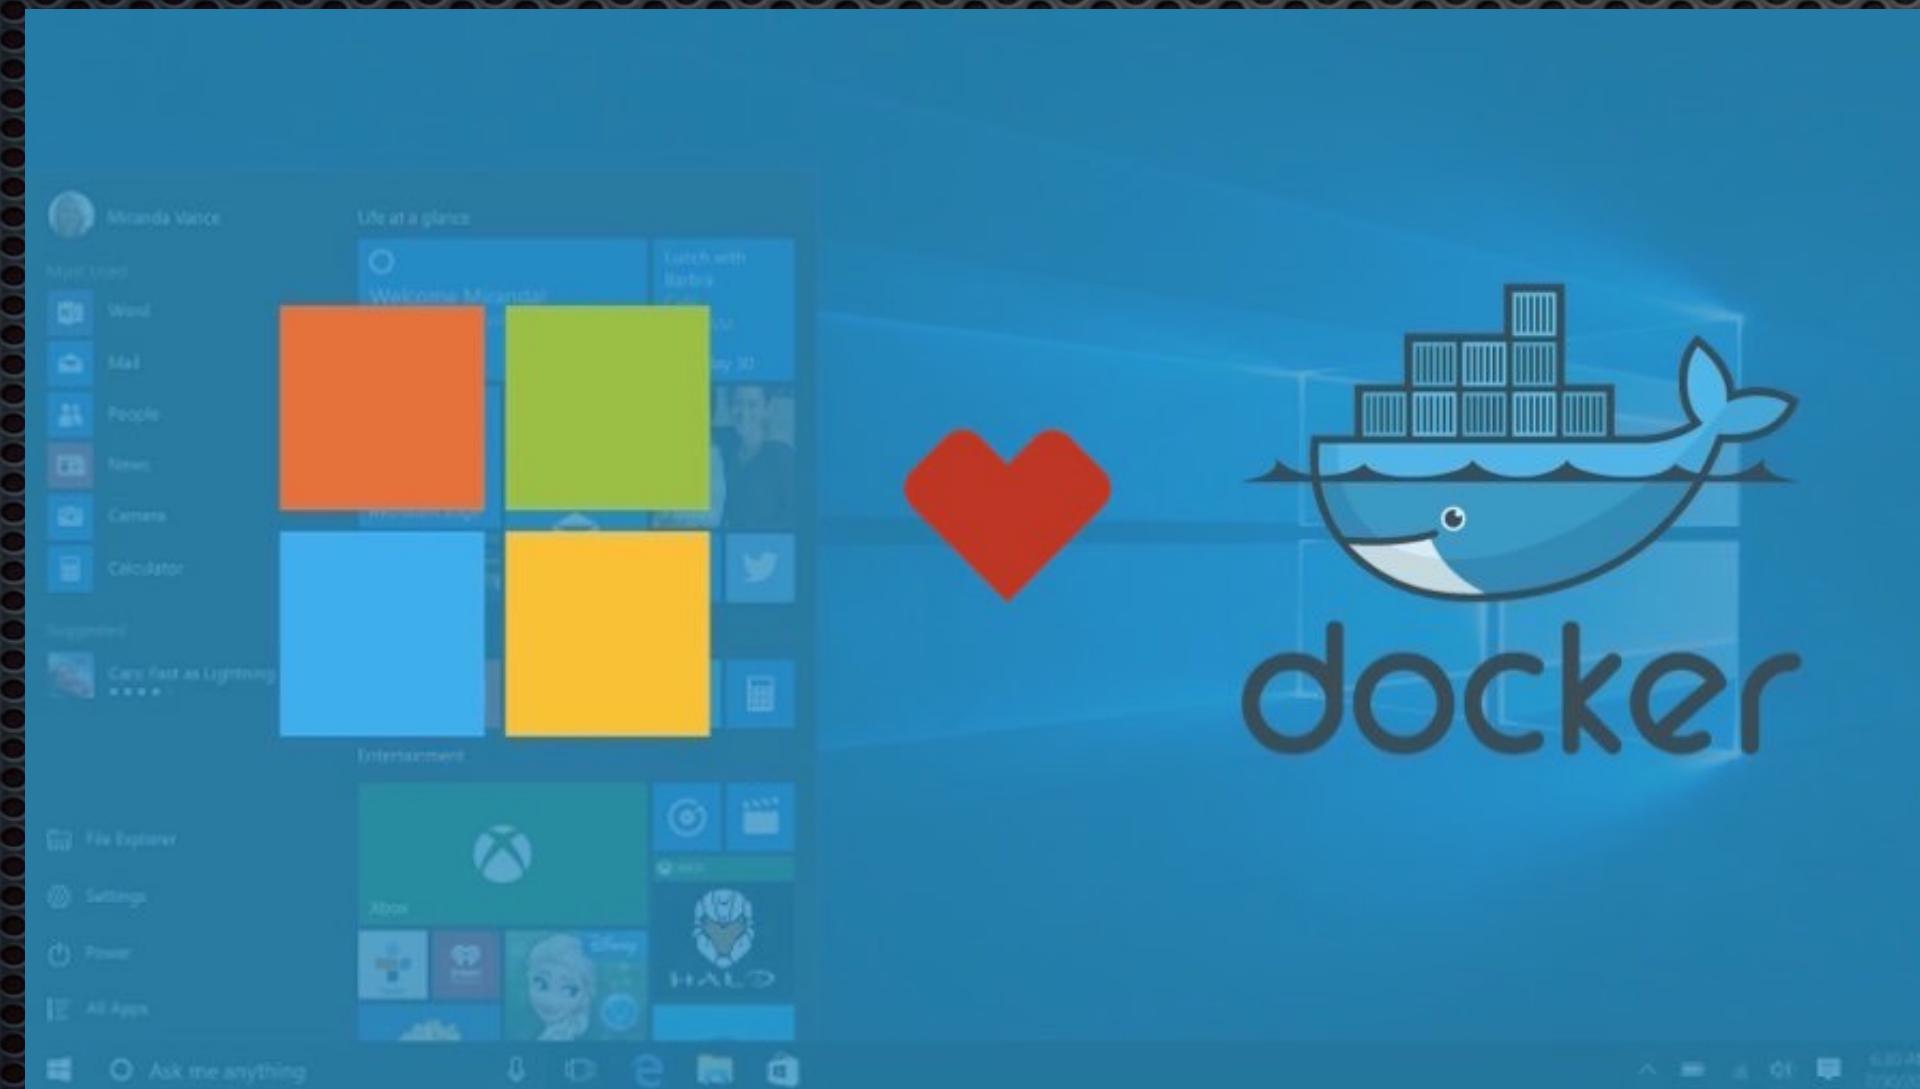
# Ahoy kubernauts

- Long time ago in a galaxy far, far away ..
- Hybrid.. what the hell?
- Providers.. why Azure?
- Tooling.. why native tools?
- Demo
- Q&A

- few++ words about me ...
- Comae Technologies - Cybersecurity startup
  - Memory-based Threat Hunting
  - No managers, no QA, no System Administrators
  - Only developers and CEO
  - memory forensics requires Windbg
- Aim: start Windows + Linux Kubernetes cluster in 15 minutes

## Why Windows?

- \$\$\$ grant from MS
- Better GPU utilization
- Office CLI tools

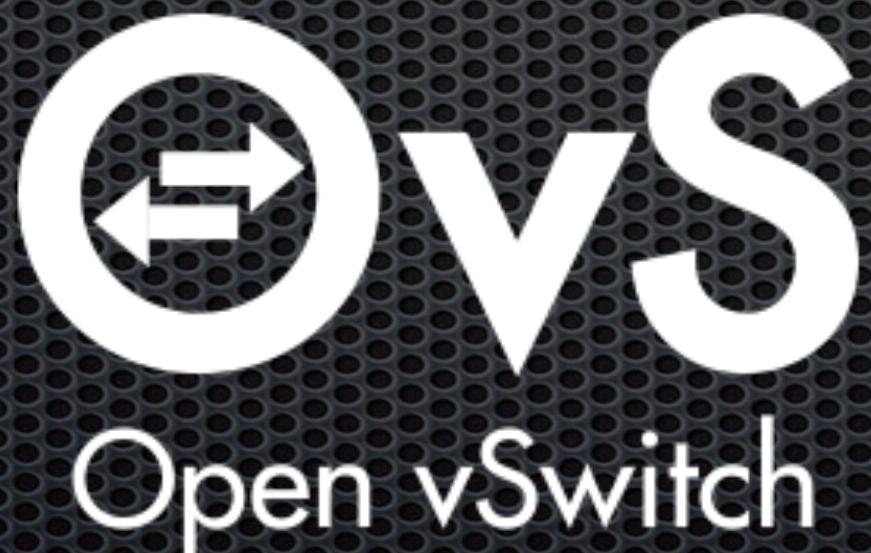


- Windows-based Utilities (windbg)
- WMV transcoding
- Religious reasons
- ...

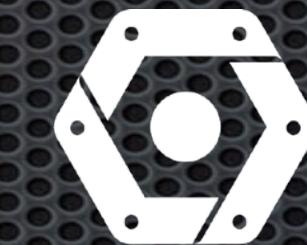


AWS - KOPs seems have issues with windows ?

GCE - mainly manual setup ?



- MS has stake in #SIG-Azure
- Works with acs-engine



- #SIG-Windows solution: roadmap
- Manual setup for now

# Provisioning

- Terraform, Vagrant, Chef
- All the above have learning curve... not best for startup
- ... anyone tried win on kube - please share experience?
- Save time and money with native cloud tools!

source ./1.0.create-azure-group-source.sh

```
export kubeGroup='kube-demo'  
export dnsPrefix='kube-demo-kiev'  
az group create --name $kubeGroup --location eastus
```

# Why native cloud tools?

- Early stage, windows containers not mature
  - Cloud resolves own tools first
  - Provided by cloud developers with inside info
    - Better architecture
    - Access to cloud resources (e.g. Primary Storage)



**But! it will lock you in to the cloud**

# acs-engine deployment

```
$ git clone https://github.com/Azure/acs-engine
$ cd ./acs-engine
$ ./scripts/devenv.sh
acs$ make bootstrap
acs$ make build
```

```
source ./1.create-acs-engine-template-json.sh
```

```
export SUBSCRIPTION_ID=`az account show | jq ".id" -r`  
export res1=`az ad sp create-for-rbac --role="Contributor" --scopes="/subscriptions/${SUBSCRIPTION_ID}"`  
export spName=$(echo $res1 | jq ".appId" -r)  
export spPassword=$(echo $res1 | jq ".password" -r)  
rm -rf acs-engine/_output  
export linuxKeyData=$(< ~/.ssh/id_rsa.pub)  
export windowsPassword='ReplaceWithSecurePassword1234'  
export dnsPrefix='my-kube-hybrid'  
mkdir acs-engine/_output
```

```
jq <acs-engine/examples/windows/kubernetes-hybrid.json  
".properties.servicePrincipalProfile.clientId=\"$spName\""  
| ".properties.servicePrincipalProfile.secret=\"$spPassword\""  
| ".properties.linuxProfile.ssh.publicKeys[0].keyData=\"$linuxKeyData\""  
| ".properties.windowsProfile.adminPassword=\"$windowsPassword\""  
| ".properties.masterProfile.dnsPrefix=\"$dnsPrefix\""" >acs-engine/_output/demo-kubernetes-hybrid.json
```

# acs-engine

- Microsoft's fork of Kubernetes
- Cluster definition [github/Azure/acs-engine](https://github.com/Azure/acs-engine) (see ./examples)
- AvailabilitySet only, VirtualMachineScaleSet not supported
- StorageAccount or ManagedDisks up to 4 disks
- Network: vnet, route table, nsg, load balancers
- Node labeling, version control... cluster upgrade?!

```
acs$ ./bin/acs-engine generate _output/demo-kubernetes-hybrid.json
```

```
source ./2.create-cluster.sh
```

```
export clusterInfo=`az group deployment create --name ${kubeGroup}-
deployment --resource-group ${kubeGroup} --template-file "./acs-engine/
_output/${dnsPrefix}/azuredploy.json" --parameters "./acs-engine/
_output/${dnsPrefix}/azuredploy.parameters.json"`

echo $clusterInfo > ./acs-engine/_output/cluster-info.json

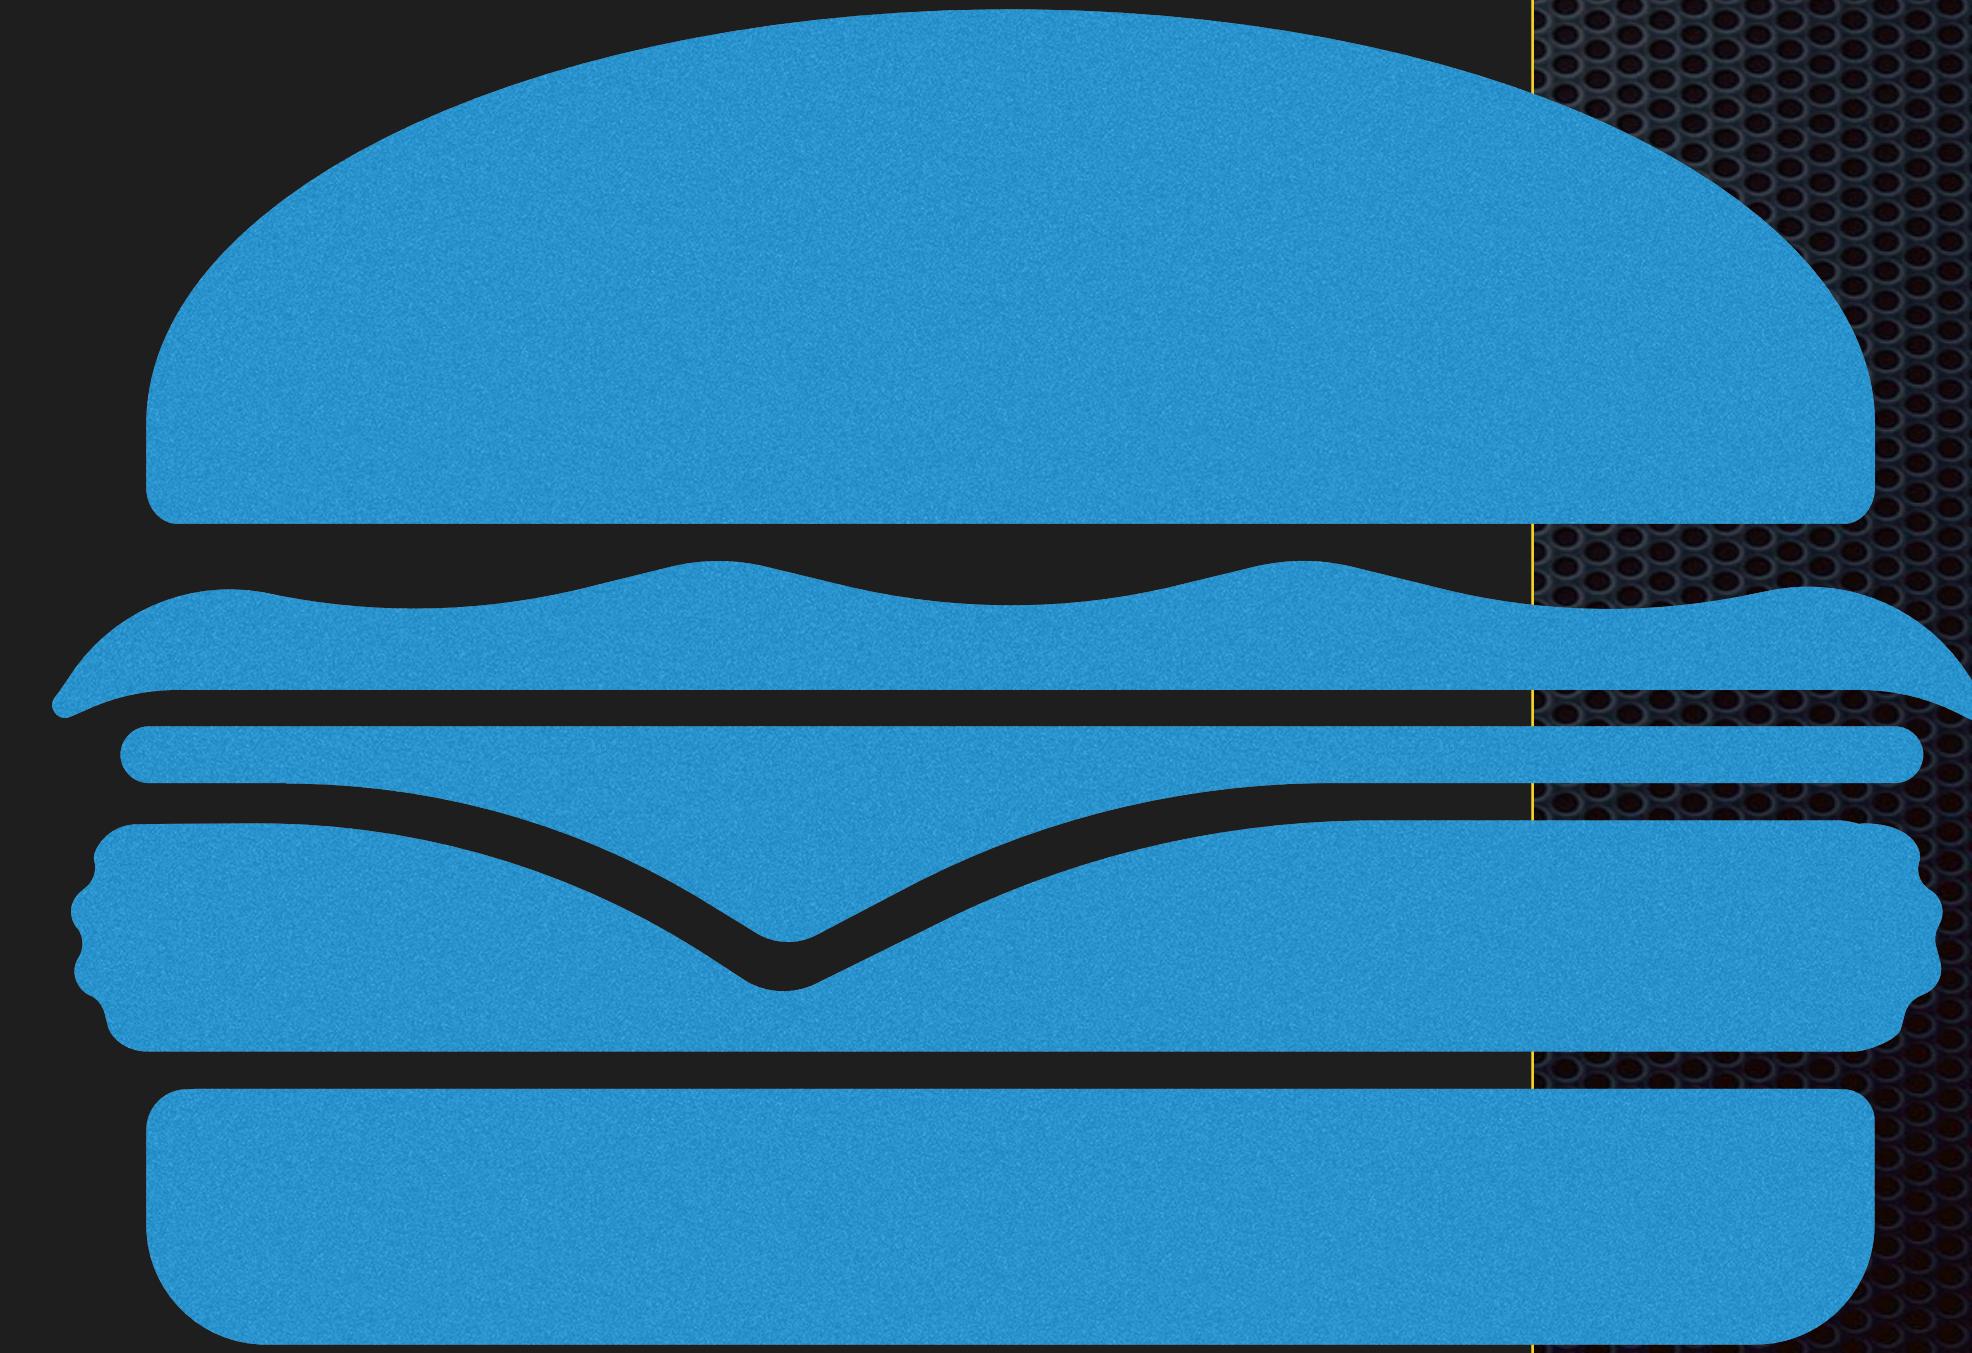
export masterHost=`echo $clusterInfo | jq
".properties.outputs.masterFQDN.value" -r`-
scp azureuser@${masterHost}:.kube/config ./acs-engine/_output/kube-config

export KUBECONFIG=$KUBECONFIG:`pwd`/acs-engine/_output/kube-config
```

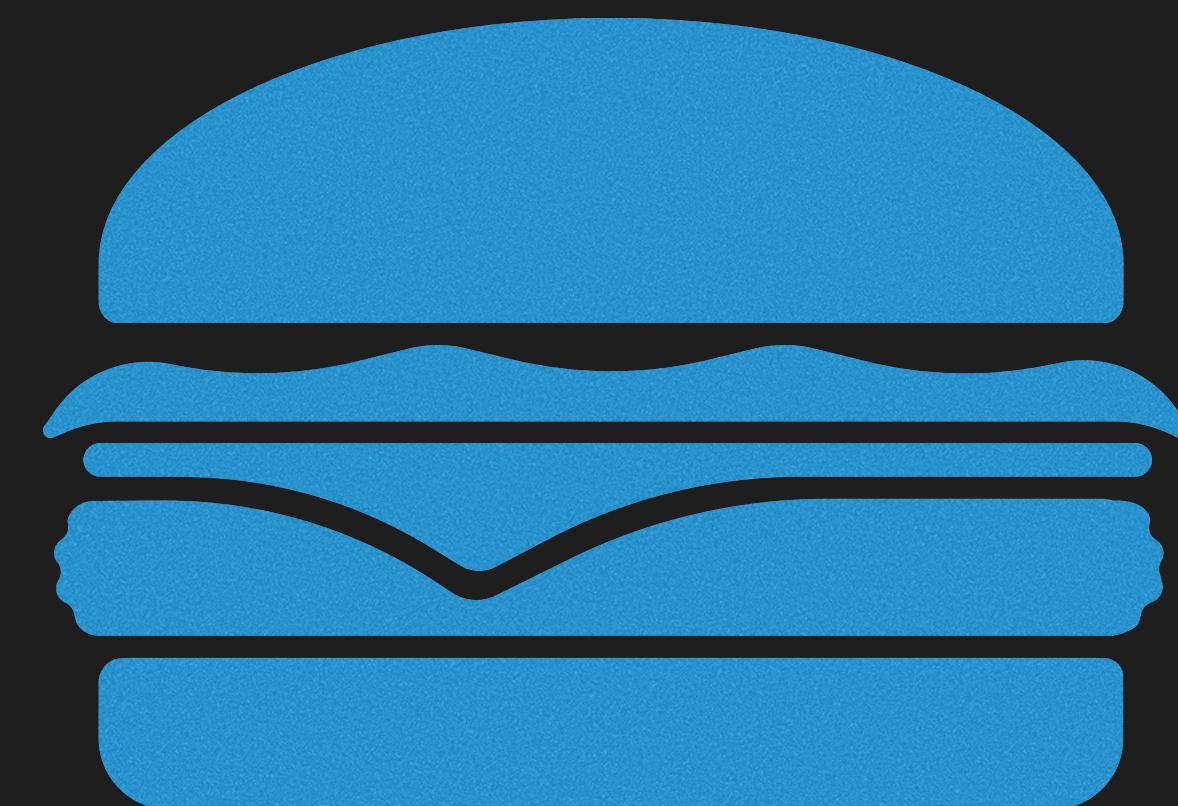
# acs-engine

```
{  
  "apiVersion": "vlabs",  
  "properties": {  
    "orchestratorProfile": {  
      "orchestratorType": "Kubernetes"  
    },  
    "masterProfile": {  
      "count": 1,  
      "dnsPrefix": "comae-kube-dev",  
      "vmSize": "Standard_D2_v2"  
    },  
    ....  
    "servicePrincipalProfile": {  
      "servicePrincipalClientID": "96352c52-ff89-4023-1911-5931234",  
      "servicePrincipalClientSecret": "5931234-ff89-4023-191196-352c52"  
    }  
  }  
}
```

```
{  
  "apiVersion": "vlabs",  
  "properties": {  
    "orchestratorProfile": {  
      "orchestratorType": "Kubernetes"  
    },  
    "masterProfile": { ... }  
    "agentPoolProfiles": [  
      {  
        "name": "linuxpool1",  
        "count": 3,  
        "vmSize": "Standard_DS3_v2",  
        "availabilityProfile": "AvailabilitySet",  
        "storageProfile": "StorageAccount",  
        "osDiskSizeGB": 30  
      },  
      {  
        "name": "windowspool2",  
        "count": 1,  
        "vmSize": "Standard_DS11_v2",  
        "availabilityProfile": "AvailabilitySet",  
        "osType": "Windows",  
        "storageProfile": "ManagedDisks",  
        "osDiskSizeGB": 256,  
        "diskSizesGB": [128, 128, 128, 128],  
      }  
    ],  
    ....  
    "servicePrincipalProfile": { ..... }  
  }  
}
```



```
{  
  "apiVersion": "vlabs",  
  "properties": {  
    "orchestratorProfile": {  
      "orchestratorType": "Kubernetes"  
    },  
    "masterProfile": { ... }  
    "agentPoolProfiles": [  
      {  
        "name": "linuxpool1",  
        ....  
      },  
      {  
        "name": "windowspool2",  
        ....  
      }  
    ],  
    "windowsProfile": {  
      "adminUsername": "azureuser",  
      "adminPassword": "notsosecret"  
    },  
    "linuxProfile": {  
      "adminUsername": "azureuser",  
      "ssh": {  
        "publicKeys": [  
          {  
            "keyData": "ssh-rsa....."  
          }  
        ]  
      }  
    },  
    "servicePrincipalProfile": { ..... }  
  }  
}
```



# Caveats (Linux)

- ssh to agent via tunnel

```
ssh azureuser@kube-demo-kiev-1.eastus.cloudapp.azure.com -L 3390:k8s-linuxpool1-30163927-0:22 -N
```

- No swap on initial linux host in Azure
  - edit /etc/waagent.conf

```
# Format if unformatted. If 'n', resource disk will not be mounted.  
ResourceDisk.Format=y  
  
# File system on the resource disk  
# Typically ext3 or ext4. FreeBSD images should use 'ufs2' here.  
ResourceDisk.Filesystem=ext4  
  
# Mount point for the resource disk  
ResourceDisk.MountPoint=/mnt  
  
# Create and use swapfile on resource disk.  
ResourceDisk.EnableSwap=y  
  
# Size of the swapfile.  
ResourceDisk.SwapSizeMB=10000
```

# Caveats (only Windows)

- Windowsservercore = 5GB minimal install. But! Nanoserver = 400MB
- One container in pod! OVS going to resolve this...
- Windows Docker volume mounted in subdir fails with fstat() \\_\(\ツ)\\_/\\_
- Workaround - mount to drive letter.
- Kubernetes volumeMount does not accept drive letter \\_\(\ツ)\\_/\\_ \\_\(\ツ)\\_/\\_
- Workaround - run **net use** from within application

# Caveats (only Windows)

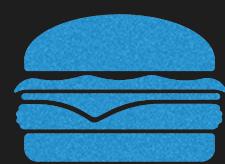
- kube-dns from Windows container sometimes fails for CNAME records
- 1500 MTU is still issue
- Windows Container caches negative results
- Windows Firewall - should not block connections from any master/agent
- RDC to agent via tunnel
- Need to know! Kubernetes agent targeting: nodeSelector vs tolerations

```
PS C:\omega> netsh interface ipv4 show subinterfaces
          MTU  MediaSenseState   Bytes In   Bytes Out  Interface
-----+-----+-----+-----+-----+
        4294967295
        1500
        1500
          1      0      0  Loopback Pseudo-Interface 3
          1  58775958  20308458 vEthernet (Container NIC 9bba4cf5)
          1 4091981579 4044400397 vEthernet (Container NIC 727da20a)
```

# Kubernetes agent targeting

By OS label:

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: gnome
spec:
  template:
    spec:
      containers:
        image: kmdgeek/node
        name: gnome
      nodeSelector:
        beta.kubernetes.io/os: windows
```

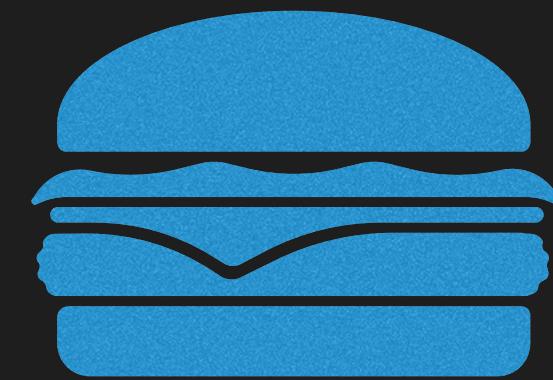


# Kubernetes agent targeting

```
$ kubectl taint node winagentnodename dedicated=windows:NoSchedule
```

Allow container on node (tolerate):

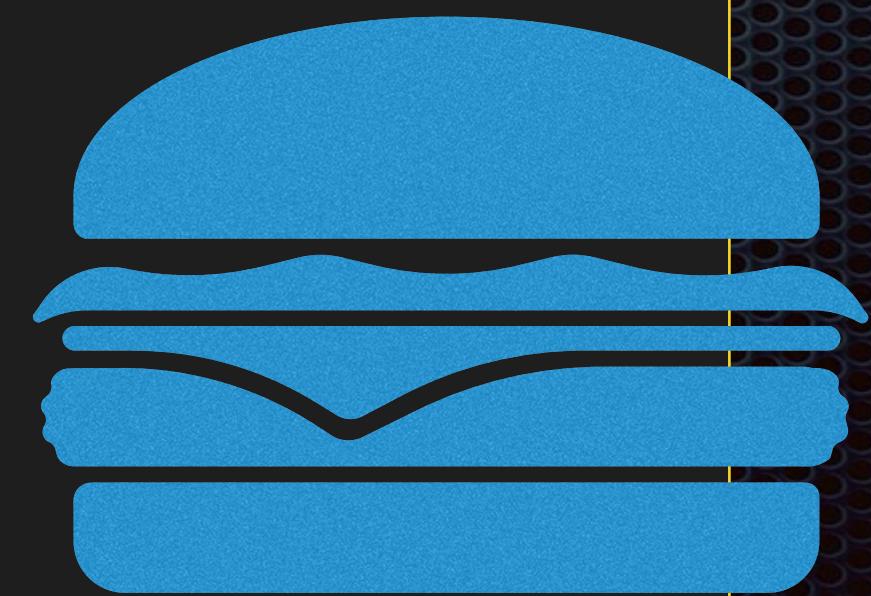
```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: gnome
spec:
  template:
    spec:
      containers:
        image: kmdgeek/node
        name: gnome
      tolerations:
      - key: "dedicated"
        operator: "Equal"
        value: "windows"
        effect: "NoSchedule"
```



# Kubernetes agent targeting

Allow on Windows and only on Windows:

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: gnome
spec:
  template:
    spec:
      containers:
        image: kmdgeek/node
        name: gnome
      nodeSelector:
        beta.kubernetes.io/os: windows
      tolerations:
      - key: "dedicated"
        operator: "Equal"
        value: "windows"
        effect: "NoSchedule"
```



# References

- <https://medium.com/@maxsparr0w/hybrid-kubernetes-in-azure-9c0c8269aca8>
- <https://github.com/dunnock/kubernetes-hybrid-in-azure>
- <https://github.com/Azure/acs-engine>
- <https://docs.microsoft.com/en-us/cli/azure/install-azure-cli?view=azure-cli-latest>
- [https://cloud.google.com/compute/docs/containers/#windows\\_known\\_issues](https://cloud.google.com/compute/docs/containers/#windows_known_issues)
- <https://kubernetes.io/docs/getting-started-guides/windows/>
- SIG Windows: <https://docs.google.com/document/d/1LWi9-NZsIZM5ITzMYoAKWXPAJwte5Ow3ySqkrGdHoLg/edit#heading=h.fz2smrec3l4>
- <https://github.com/kubernetes/kubernetes/issues/38999>
- <https://github.com/kubernetes/kubernetes/issues/48638>
- <https://github.com/nodejs/node/issues/8897>
- <https://github.com/kubernetes/kubernetes/issues/49398>